BrainStation Capstone Sprint 2

Ethan Alwaise

April 2, 2024

The Stack Exchange

A network of Q&A websites for various fields

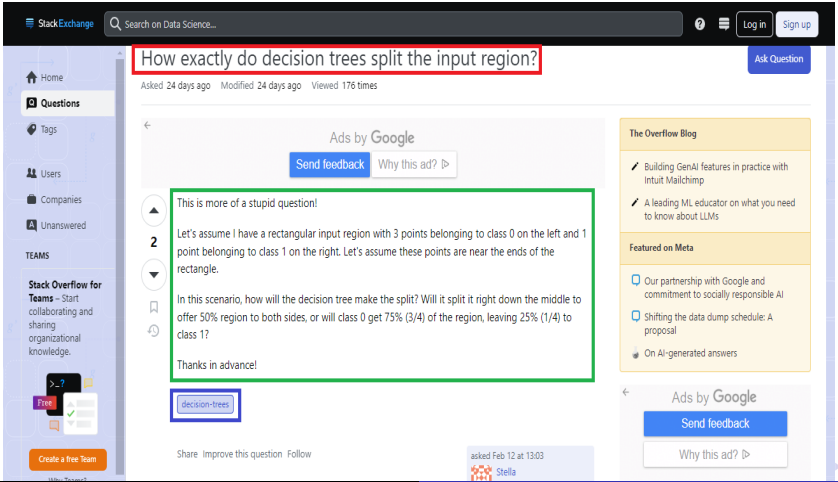

The Problem

1. Questions:
    - What factors affect the likelihood that a question is answered?
    - What patterns can we observe in posting activity over time?
2. Who should care?
    - Users who want to optimize their questions.
    - Stakeholders in the Stack Exchange.

The Problem

- Can we predict if a given question will be answered within a week?
- Can we predict the number of posts made on a given day?

The Data

	PostTypeId	CreationDate	Score	ViewCount	Body	LastActivityDate	Title	Tags	AnswerCount	CommentCount	LastEditDate
0	1	2014-05-13T23:58:30.457	9	959.0	<p>I've always been interested in machine lear...	2014-05-14T00:36:31.077	How can I do simple machine learning without h...	<machine-learning>	1.0	1	None
1	1	2014-05-14T00:11:06.457	4	503.0	<p>As a researcher and instructor, I'm looking...	2014-05-16T13:45:00.237	What open-source books (or other materials) pr...	<education> <open-source>	3.0	4	2014-05-16T13:45:00.237
2	2	2014-05-14T00:36:31.077	5	NaN	<p>Not sure if this fits the scope of this SE,...	2014-05-14T00:36:31.077	None	None	NaN	0	None
3	2	2014-05-14T00:53:43.273	13	NaN	<p>One book that's freely available is "The El...	2014-05-14T00:53:43.273	None	None	NaN	1	None
4	1	2014-05-14T01:25:59.677	26	1925.0	<p>I am sure data science as will be discussed...	2020-08-16T13:01:33.543	Is Data Science the Same as Data Mining?	<data-mining> <definitions>	4.0	1	2014-06-17T16:17:20.473

EDA (Answer Prediction)

Class balance:

$\sim 65\%$ of questions are answered within 7 days.

EDA (Daily Post Activity)

We focus on 2020 onward since we observe two opposite trends:

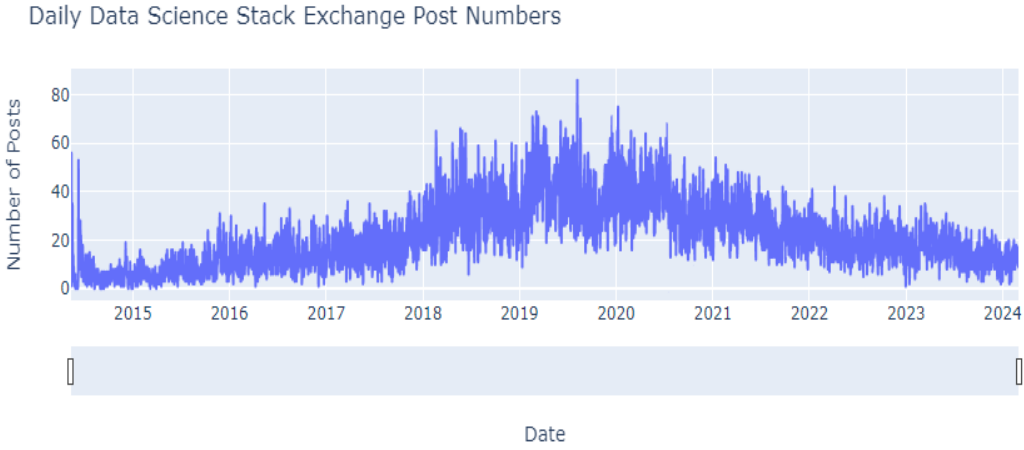

Answer Prediction (Logistic Regression)

Feature Engineering and Data Processing

Steps:

- Count math equations, lines of code, etc. in questions.
- Create dummy variables for question tags.
- Vectorize the text (bag of words).

Answer Prediction (Logistic Regression)

Model Performance

		Precision	Recall	F1-score
TRAIN	0 (Unanswered)	0.71	0.39	0.50
	1 (Answered)	0.74	0.92	0.82
	Accuracy			0.74
TEST	0 (Unanswered)	0.49	0.19	0.27
	1 (Answered)	0.67	0.90	0.77
	Accuracy			0.65

Answer Prediction (Logistic Regression)

Next Steps

- Improve tokenization (include non-alphabetic characters)
- More sophisticated text processing.
- Try dimensionality reduction and hyperparameter optimization.
- Experiment with other types of models (random forest, neural nets, etc.).

Daily Post Activity Prediction (SARIMA)

EDA and Hyperparameter Selection

Modeling steps:

- Count number of posts per day and observe trends.
- Determine order of differencing using unit root testing.
- Look at partial autocorrelations to determine autoregressive order.

Daily Post Activity Prediction (SARIMA)

Model Predictions

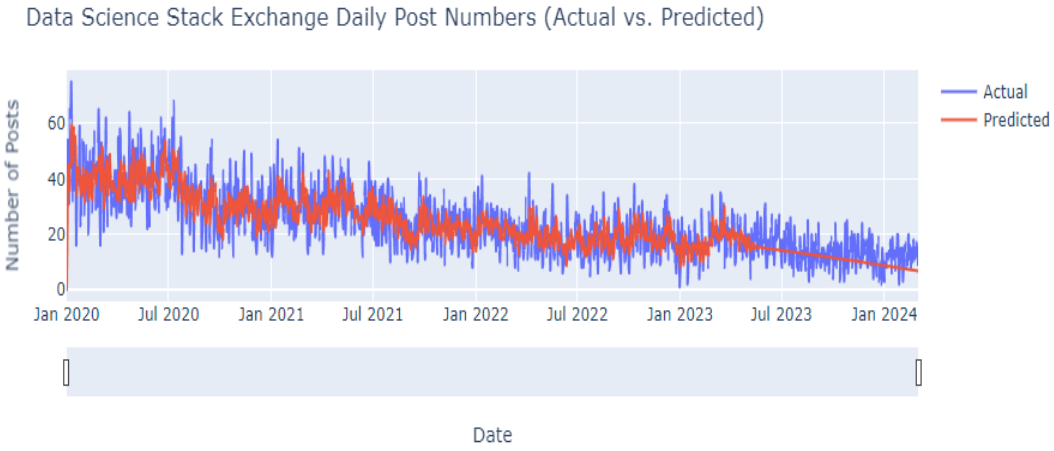

Daily Post Activity Prediction (SARIMA)

Model Performance

	MAPE	RMSE
TRAIN	30.0%	7.96
TEST	40.0%	5.72

Daily Post Activity Prediction (SARIMA)

Next Steps

- Tune hyperparameters.
- Incorporate seasonality.
- Experiment with rolling averages / monthly numbers.